

CSCI 3202: Intro to Artificial Intelligence

Lecture 8: Optimality

Rachel Cox
Department of
Computer Science



Review

Uniform-cost search:

$$f(n) = g(n) \quad (\text{cost to get to } n)$$

Greedy:

$$f(n) = h(n) \quad (\text{estimated cost to get from } n \text{ to goal})$$

A*:

$$f(n) = g(n) + h(n)$$

(estimated total cost of cheapest solution through n)

Optimality

Conditions for Optimality: Admissibility & Consistency

- $h(n)$ must be **admissible** - an admissible heuristic is one that never overestimates the cost to reach the goal.
- $h(n)$ is **consistent** if, for every node n and every successor n' of n generated by any action a , the estimated cost of reaching the goal from n is no greater than the step cost of getting to n' plus the estimated cost of reaching the goal from n' :

$$h(n) \leq c(n, a, n') + h(n')$$

Optimality of A* Search

A* is **optimally efficient** for any given heuristic: No other optimal algorithm is guaranteed to expand fewer nodes than A*

- Recall: A* expands all nodes with $f(n) < C^*$, where C^* is the cost of the optimal solution path.
- Any algorithm that does not expand all nodes with $f(n) < C^*$ risks missing a better solution path.

Optimality of A* Search

A* (graph) is optimal if the heuristic $h(n)$ is consistent.

Based on two key facts:

1. If $h(n)$ is consistent, then the values of $f(n)$ along any path are nondecreasing.
2. Whenever A* selects a node n for expansion, the optimal path to that node has been found.

Optimality of A* Search

If $h(n)$ is consistent, then the values of $f(n)$ along any path are nondecreasing.

Example: Prove the above statement.

Optimality of A* Search

Whenever A* selects a node n for expansion, the optimal path to that node has been found.

Example: Prove the above statement.

Optimality of A* Search

A* (graph) is optimal if the heuristic $h(n)$ is consistent.

Based on two key facts:

1. If $h(n)$ is consistent, then the values of $f(n)$ along any path are nondecreasing.
2. Whenever A* selects a node n for expansion, the optimal path to that node has been found.

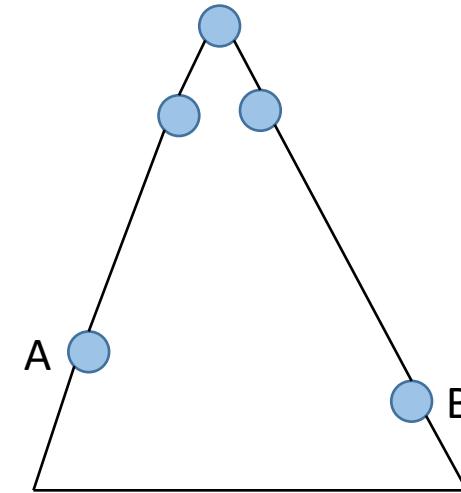
➤ **So the first goal node to be expanded took the lowest-cost path, and all later goal node expansions are at least as expensive.**

Optimality of A* Search

Assume:

- A is an optimal Goal Node
- B is a suboptimal Goal node
- h is admissible

Claim: A will exit the frontier before B.



Optimality of A* Search

So A* is **optimal**, **complete**, and **optimally efficient**.

Why do we even care about other search algorithms?

- **Number of nodes** to expand along the goal contour is still **exponential** in depth of solution/length of solution path.
- Absolute error: $\Delta := h^* - h$
 - h^* = actual cost from root to goal
 - h = heuristic you used
- Relative error: $\epsilon := (h^* - h)/h^*$

A* Search

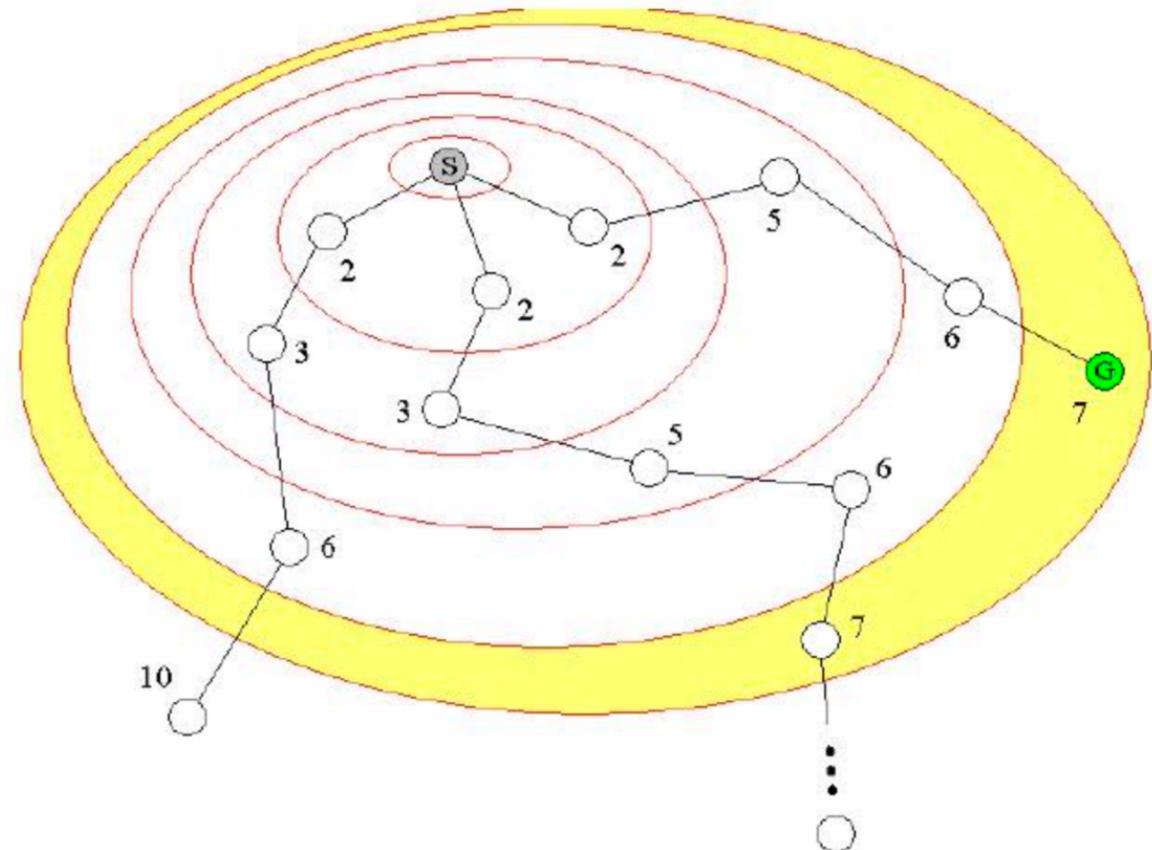
Complexity depends strongly on state space characterization

- Single goal, tree, reversible actions $\rightarrow O(b^\Delta)$, or $O(b^{\epsilon d})$ with constant step costs (d is solution depth)
 - Δ typically is proportional to the path cost h^* , so ϵ is pretty much constant (or growing with d), and we can rewrite: $O((b^\epsilon)^d)$
 - The effective branching factor is really b^ϵ .
 - Important to choose as good of a heuristic as we can.
- Many goal states/near-goal states can be a problem -- need to expand a *lot* of branches.

A* Search: alternatives

Space complexity can be a burden for A*

- Keeps all generated nodes in memory
- Iterative Deepening A* (IDA*)
 - Just like standard IDS, but instead of extending the search depth, extend the allowed f cost.
 - Search out to a particular contour



A* Search: alternatives

- **Recursive best-first search (RBFS)**
 - Like standard DFS, but keeps track of the best alternative path's f -cost
 - Once the path you're digging down exceeds the best alternative path, switch over to the back-up path
 - As RBFS back-tracks, each node along the back-tracked path it replaces the f -value with that of the cheapest child node.
 - Remembers the best leaf in the sub-tree, so RBFS knows whether it's worth it to go back down that road later.
 - Still has problems of indecision
 - Expanding a new path makes the unexpanded alternatives look better.

Recursive best-first Search (RBFS)

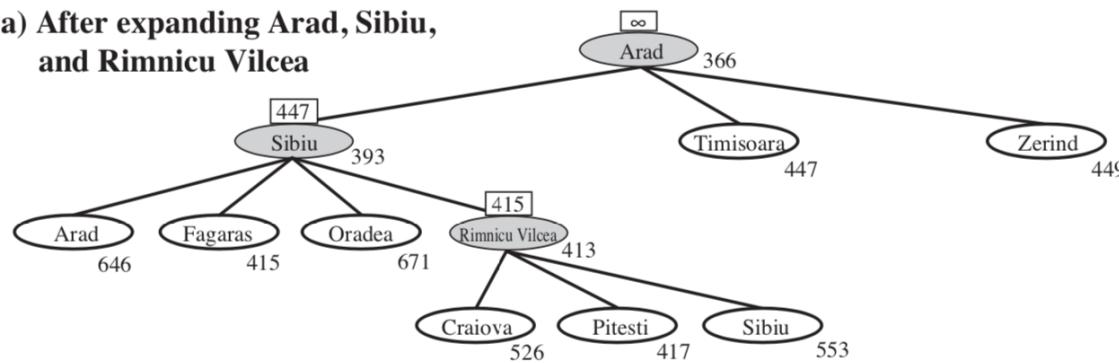
```
function RECURSIVE-BEST-FIRST-SEARCH(problem) returns a solution, or failure
    return RBFS(problem, MAKE-NODE(problem.INITIAL-STATE),  $\infty$ )

function RBFS(problem, node, f-limit) returns a solution, or failure and a new f-cost limit
    if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)
    successors  $\leftarrow$  []
    for each action in problem.ACTIONS(node.STATE) do
        add CHILD-NODE(problem, node, action) into successors
    if successors is empty then return failure,  $\infty$ 
    for each s in successors do /* update f with value from previous search, if any */
        s.f  $\leftarrow$  max(s.g + s.h, node.f)
    loop do
        best  $\leftarrow$  the lowest f-value node in successors
        if best.f  $>$  f-limit then return failure, best.f
        alternative  $\leftarrow$  the second-lowest f-value among successors
        result, best.f  $\leftarrow$  RBFS(problem, best, min(f-limit, alternative))
        if result  $\neq$  failure then return result
```

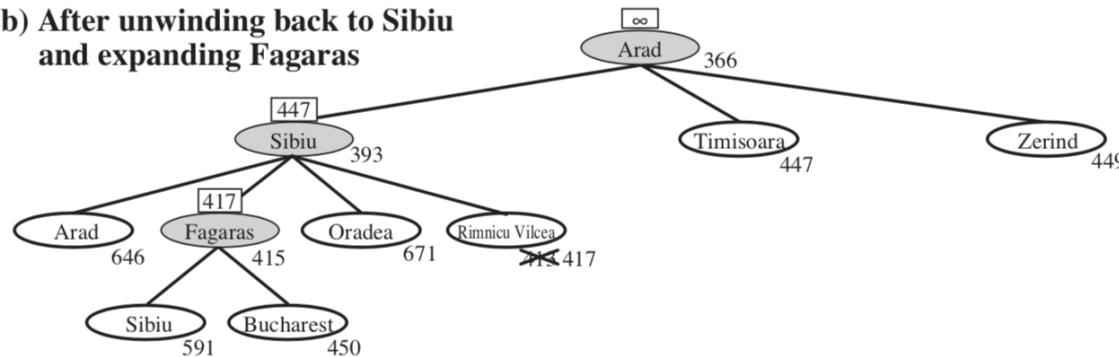
Figure 3.26 The algorithm for recursive best-first search.

Recursive best-first Search (RBFS)

(a) After expanding Arad, Sibiu, and Rimnicu Vilcea



(b) After unwinding back to Sibiu and expanding Fagaras



(c) After switching back to Rimnicu Vilcea and expanding Pitesti

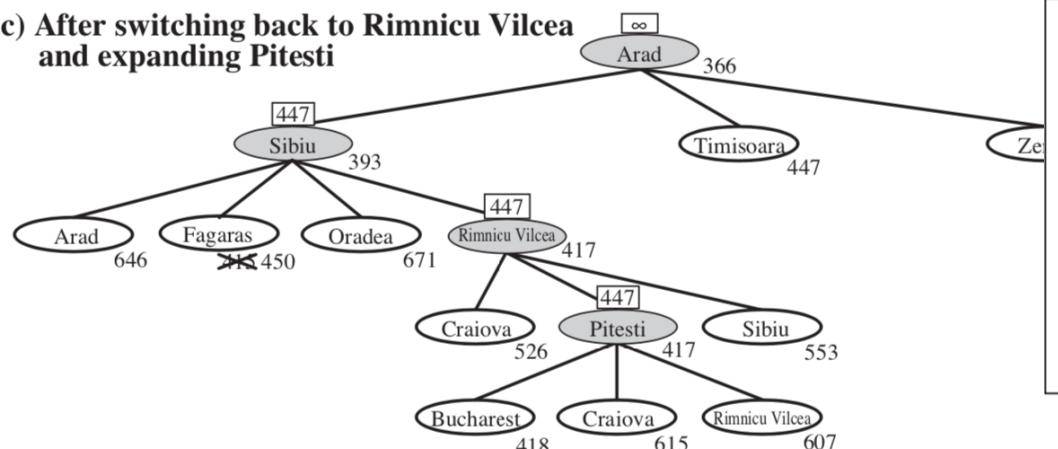


Figure 3.27 Stages in an RBFS search for the shortest route to Bucharest. The f -limit value for each recursive call is shown on top of each current node, and every node is labeled with its f -cost. (a) The path via Rimnicu Vilcea is followed until the current best leaf (Pitesti) has a value that is worse than the best alternative path (Fagaras). (b) The recursion unwinds and the best leaf value of the forgotten subtree (417) is backed up to Rimnicu Vilcea; then Fagaras is expanded, revealing a best leaf value of 450. (c) The recursion unwinds and the best leaf value of the forgotten subtree (450) is backed up to Fagaras; then Rimnicu Vilcea is expanded. This time, because the best alternative path (through Timisoara) costs at least 447, the expansion continues to Bucharest.

A* Search: alternatives

- IDA* and RBFS use very little memory
 - For example:
 - Between iterations, IDA* keeps only the current f -cost limit
 - RBFS has memory benefits of DFS, but at the cost of potential time inefficiency
 - Those who don't remember the past, are doomed to repeat it (maybe).

A* Search: alternatives

Memory-bounded A* and Simplified Memory-bounded A* (MA* and SMA*)

- SMA* -- general notes and subtleties
 - Expands the best leaf until memory is full
 - Then expands the best leaf and deletes the worst
 - What if all the leaves have the same f -value?
 - Then expand the newest leaf and delete the oldest
 - What if there is only one leaf?
 - That's ok -- then there is a single solution path from root to goal, and it's taking up all available memory, so we've failed anyway

Next Time

Heuristics